

# M5 Forecasting accuracy

The objective of this notebook consists in forecasting unit sales for given products and stores.

The notebook is divided into the following parts:

- **Preprocessing**: data preprocessing to make them more suitable to the final goal;
- **Exploratory Data Analysis (EDA)**: data exploration, looking for possible patterns or correlations;
- **Feature Engineering (FE)**: computation of new features starting from available ones;
- **Model implementation**: implementation of a model for each store;
- **Predictions**: predictions making for public and private data;
- **Submissions**: computation of partial and final submissions;
- **Evaluation**: error computation and its decomposition for each aggregation level.

```
In [ ]: # Run the following three cells only if using colab
```

```
In [ ]: !pip install lightgbm==3.3.2
```

```
In [1]: !git clone https://github.com/michelevece/vece-m5-forecasting-accuracy
```

```
In [2]: %cd vece-m5-forecasting-accuracy/data/code
```

## 0) Libraries

```
In [1]: # libraries
import os
import zipfile
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import tqdm
import lightgbm as lgbm
from sklearn.metrics import mean_squared_error as mse

# my modules
import m5_utils as utils
import m5_preprocessing as preprocessing
import m5_fe as fe
import m5_wrmsse_evaluator as wrmsse_evaluator

# target column
TARGET = 'sales'

# first day of the private leaderboard
D_PUBLIC = 1914
# first day of the public leaderboard
D_PRIVATE = 1942
```

## 1) Preprocessing

Data concerns three main aspects:

- Prices

- Calendar
- Sales

In the following:

- data is extracted;
- data undergo downcasting (in order to reduce the memory storage);
- some columns are added or removed;
- data is stored in parquet format.

```
In [3]: utils.extract_data()

data extracted
```

```
In [4]: calendar, prices, sales = utils.read_data_csv()
```

## Prices

The prices dataframe contains the price ( `sell_price` ) at which a product ( `item_id` ) is sold in a store ( `store_id` ) in a given week of a given year ( `wm_yr_wk` ).

In this dataframe, weeks start on Saturday and end on Friday.

```
In [5]: prices.head()
```

```
Out[5]:
```

	store_id	item_id	wm_yr_wk	sell_price
0	CA_1	HOBBIES_1_001	11325	9.58
1	CA_1	HOBBIES_1_001	11326	9.58
2	CA_1	HOBBIES_1_001	11327	8.26
3	CA_1	HOBBIES_1_001	11328	8.26
4	CA_1	HOBBIES_1_001	11329	8.26

```
In [6]: prices.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 6841121 entries, 0 to 6841120
Data columns (total 4 columns):
#   Column      Dtype
---  -
0   store_id    object
1   item_id     object
2   wm_yr_wk    int64
3   sell_price  float64
dtypes: float64(1), int64(1), object(2)
memory usage: 208.8+ MB
```

```
In [7]: # casting
prices['store_id'] = prices['store_id'].astype('category')
prices['item_id'] = prices['item_id'].astype('category')

prices['wm_yr_wk'] = prices['wm_yr_wk'].astype('int16')
prices['sell_price'] = prices['sell_price'].astype('float32')
```

```
In [8]: # missing values check
print('Missing values:', prices.isna().any().any())
```

```
print('min week:', prices['wm_yr_wk'].min())
print('max week:', prices['wm_yr_wk'].max())
```

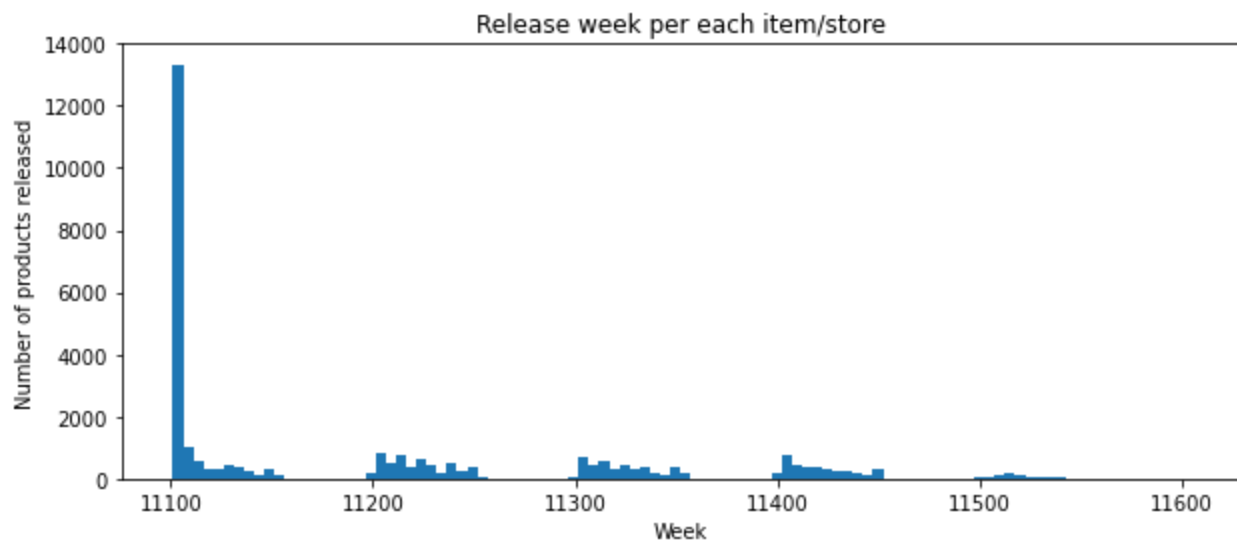
```
Missing values: False
min week: 11101
max week: 11621
```

## Release

Not all products are released in the same period. However, most of them are available since the first year

```
In [9]: plt.figure(figsize=(10,4))
releases = prices.groupby(['store_id', 'item_id'])['wm_yr_wk'].min()
plt.hist(releases, bins=100)
plt.xlabel('Week')
plt.ylabel('Number of products released')
plt.title('Release week per each item/store')
```

```
Out[9]: Text(0.5, 1.0, 'Release week per each item/store')
```



## Calendar

Calendar contains information about 1969 days from 29-01-2011 to 19-06-2016 .

In details, there are:

- `wm_yr_wk` : a code that identifies a week of a given year;
- `weekday` , `wday` : day of the week;
- `month` , `year` : month and year;
- `d` : day (1-1969);
- `event_name_N` , `event_type_N` : name and category of an event;
- `snap_XX` : presence of SNAP promotion in the state `XX` ;

```
In [10]: calendar.head()
```

Out[10]:

	date	wm_yr_wk	weekday	wday	month	year	d	event_name_1	event_type_1	event_name_2	event_ty
0	2011-01-29	11101	Saturday	1	1	2011	d_1	NaN	NaN	NaN	
1	2011-01-30	11101	Sunday	2	1	2011	d_2	NaN	NaN	NaN	
2	2011-01-31	11101	Monday	3	1	2011	d_3	NaN	NaN	NaN	
3	2011-02-01	11101	Tuesday	4	2	2011	d_4	NaN	NaN	NaN	
4	2011-02-02	11101	Wednesday	5	2	2011	d_5	NaN	NaN	NaN	

In [11]:

calendar.info()

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1969 entries, 0 to 1968
Data columns (total 14 columns):
#   Column          Non-Null Count  Dtype
---  -
0   date            1969 non-null   object
1   wm_yr_wk        1969 non-null   int64
2   weekday         1969 non-null   object
3   wday            1969 non-null   int64
4   month           1969 non-null   int64
5   year            1969 non-null   int64
6   d               1969 non-null   object
7   event_name_1    162 non-null    object
8   event_type_1    162 non-null    object
9   event_name_2    5 non-null      object
10  event_type_2     5 non-null      object
11  snap_CA         1969 non-null   int64
12  snap_TX         1969 non-null   int64
13  snap_WI         1969 non-null   int64
dtypes: int64(7), object(7)
memory usage: 215.5+ KB
```

In [12]:

```
# casting
calendar['date'] = calendar['date'].astype('datetime64')

calendar['wm_yr_wk'] = calendar['wm_yr_wk'].astype('int16')
calendar['month'] = calendar['month'].astype('int8')
calendar['year'] = calendar['year'].astype('int16')

# delete 'd_' prefix
calendar['d'] = calendar['d'].apply(lambda x: x[2:]).astype('int16')

calendar['snap_CA'] = calendar['snap_CA'].astype('int8')
calendar['snap_TX'] = calendar['snap_TX'].astype('int8')
calendar['snap_WI'] = calendar['snap_WI'].astype('int8')
```

In [13]:

calendar.describe(datetime\_is\_numeric=True)

Out[13]:

	date	wm_yr_wk	wday	month	year	d	snap_CA	snap_TX	
count	1969	1969.000000	1969.000000	1969.000000	1969.000000	1969.000000	1969.000000	1969.000000	196
mean	2013-10-09 00:00:00	11347.086338	3.997461	6.325546	2013.288471	985.000000	0.330117	0.330117	
min	2011-01-29 00:00:00	11101.000000	1.000000	1.000000	2011.000000	1.000000	0.000000	0.000000	
25%	2012-06-04 00:00:00	11219.000000	2.000000	3.000000	2012.000000	493.000000	0.000000	0.000000	
50%	2013-10-09 00:00:00	11337.000000	4.000000	6.000000	2013.000000	985.000000	0.000000	0.000000	
75%	2015-02-13 00:00:00	11502.000000	6.000000	9.000000	2015.000000	1477.000000	1.000000	1.000000	
max	2016-06-19 00:00:00	11621.000000	7.000000	12.000000	2016.000000	1969.000000	1.000000	1.000000	
std	NaN	155.277043	2.001141	3.416864	1.580198	568.545659	0.470374	0.470374	

Events

Some day have more than one event. Since these days are only 5, event\_name\_2 and event\_type\_2 are removed

In [14]:

calendar[calendar.columns[7:11]].describe()

Out[14]:

	event_name_1	event_type_1	event_name_2	event_type_2
count	162	162	5	5
unique	30	4	4	2
top	SuperBowl	Religious	Father's day	Cultural
freq	6	55	2	4

In [15]:

calendar.drop(columns=['event\_name\_2', 'event\_type\_2'], inplace=True)

In [16]:

```
# fill empty event records
calendar['event_name_1'].fillna('no_event', inplace=True)
calendar['event_type_1'].fillna('no_event', inplace=True)

calendar['event_name_1'] = calendar['event_name_1'].astype('category').astype('int8')
calendar['event_type_1'] = calendar['event_type_1'].astype('category').astype('int8')
```

new columns

In [17]:

```
# day of the week, from 0 (monday) to 6 (sunday)
calendar.insert(3, 'dayofweek', calendar['date'].dt.dayofweek.astype('int8'))
# day of the month, 1-31
calendar.insert(4, 'dayofmonth', calendar['date'].dt.day.astype('int8'))
# day of the year, 1-366
```

```
calendar.insert(5, 'dayofyear', calendar['date'].dt.dayofyear.astype('int16'))
# week in the month, 1-5
calendar.insert(2, 'weekofmonth', calendar['dayofmonth'].apply(lambda x: (x - 1) // 7 + 1))
# week in the year, 1-53
calendar.insert(3, 'weekofyear', (calendar['wm_yr_wk'] % 100).astype('int8'))
```

```
In [18]: # align non-bisestile year to bisestile year
idx = calendar[(calendar['year'].isin([2011,2013,2014,2015])) & (calendar['dayofyear']>3)
calendar.loc[idx, 'dayofyear'] +=1
```

```
In [19]: # check no missing values in the 'date' column
(calendar['date'][1:].reset_index(drop=True) - calendar['date'][:-1]).value_counts()
```

```
Out[19]: 1 days      1968
Name: date, dtype: int64
```

```
In [20]: print('date min:', calendar['date'].min().date())
print('date max:', calendar['date'].max().date())
```

```
date min: 2011-01-29
date max: 2016-06-19
```

```
In [21]: calendar.drop(columns=['weekday', 'wday', 'date'], inplace=True)
calendar.set_index('d', inplace=True)
```

```
In [22]: calendar
```

```
Out[22]:
```

	wm_yr_wk	weekofmonth	weekofyear	dayofweek	dayofmonth	dayofyear	month	year	event_name_1
<b>d</b>									
<b>1</b>	11101	5	1	5	29	29	1	2011	no_event
<b>2</b>	11101	5	1	6	30	30	1	2011	no_event
<b>3</b>	11101	5	1	0	31	31	1	2011	no_event
<b>4</b>	11101	1	1	1	1	32	2	2011	no_event
<b>5</b>	11101	1	1	2	2	33	2	2011	no_event
...	...	...	...	...	...	...	...	...	...
<b>1965</b>	11620	3	20	2	15	167	6	2016	no_event
<b>1966</b>	11620	3	20	3	16	168	6	2016	no_event
<b>1967</b>	11620	3	20	4	17	169	6	2016	no_event
<b>1968</b>	11621	3	21	5	18	170	6	2016	no_event
<b>1969</b>	11621	3	21	6	19	171	6	2016	NBAFinalsEnd

1969 rows × 13 columns

## Sales

In the directory there are two datasets:

- `sales_train_validation` : historical daily unit sales per product and store in  $[d_1; d_{1913}]$ , in which the last 28 days represent the **input** for the **public** leaderboard;
- `sales_train_evaluation` : historical daily unit sales per product and store in  $[d_1; d_{1941}]$ , in which the last 28 days represent both:

- the **label** for the **public** leaderboard and
- the **input** for the **private** leaderboard.

In the following, only the last dataframe is used.

```
In [23]: sales.head()
```

```
Out[23]:
```

	id	item_id	dept_id	cat_id	store_id	state_id	d_1	d_2	d_3	d_4	...
0	HOBBIES_1_001_CA_1_evaluation	HOBBIES_1_001	HOBBIES_1	HOBBIES	CA_1	CA	0	0	0	0	...
1	HOBBIES_1_002_CA_1_evaluation	HOBBIES_1_002	HOBBIES_1	HOBBIES	CA_1	CA	0	0	0	0	...
2	HOBBIES_1_003_CA_1_evaluation	HOBBIES_1_003	HOBBIES_1	HOBBIES	CA_1	CA	0	0	0	0	...
3	HOBBIES_1_004_CA_1_evaluation	HOBBIES_1_004	HOBBIES_1	HOBBIES	CA_1	CA	0	0	0	0	...
4	HOBBIES_1_005_CA_1_evaluation	HOBBIES_1_005	HOBBIES_1	HOBBIES	CA_1	CA	0	0	0	0	...

5 rows × 1947 columns

```
In [24]: sales.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 30490 entries, 0 to 30489
Columns: 1947 entries, id to d_1941
dtypes: int64(1941), object(6)
memory usage: 452.9+ MB
```

```
In [25]: # rename day columns by deleting 'd_'
sales.rename(columns={x: x[2:] for x in sales.columns[6:]}, inplace=True)

# casting
cols = sales.columns.tolist()
sales_1 = sales[cols[:6]].astype('category')
sales_2 = sales[cols[6:]].astype('int16')
sales = pd.concat([sales_1, sales_2], axis=1)

# delete 'evaluation' suffix
sales['id'] = sales['id'].apply(lambda x: x[:-11])

del sales_1, sales_2, cols
```

```
In [26]: sales.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 30490 entries, 0 to 30489
Columns: 1947 entries, id to 1941
dtypes: category(6), int16(1941)
memory usage: 114.4 MB
```

## Write on disk

```
In [27]: dst = '../processed/'

if not os.path.exists(dst):
    os.makedirs(dst)

# save data
prices.to_parquet(dst + '/m5_prices.pqt')
calendar.to_parquet(dst + '/m5_calendar.pqt')
sales.to_parquet(dst + '/m5_sales.pqt')
```

## 2) Data exploration

```
In [28]: calendar, prices, sales = utils.read_data_pqt()
```

```
In [29]: %%time
# remove zeros sales before release date
df = preprocessing.remove_leading_zeros(sales, calendar)
```

CPU times: total: 37.2 s

Wall time: 37.4 s

### Unit sales per state, store, category and departments

It is possible to notice that the unit sales are **not** uniformly distributed among states, store, category and departments.

- Sales in CA represent more than 40% of the total.
- In California there is both the store with the highest amount of sales ( CA\_3 ) and the store with the lowest amount ( CA\_1 ).
- About 70% of the total sales comes from FOODS products, of which FOODS\_3 represents the major part.

```
In [30]: fig, ax = plt.subplots(1, 4, figsize=(18, 4))

# flatten
#ax = [cell for row in ax for cell in row]


df.groupby('state_id')['sales'].sum().plot.pie(ax=ax[0], autopct='%d%%')
ax[0].set_title('Unit sales per state')
ax[0].set_ylabel('')

df.groupby('store_id')['sales'].sum().sort_values().plot.barh(ax=ax[1]) #barh(figsize=(6
ax[1].set_title('Unit sales per store')
ax[1].set_ylabel('stores')
ax[1].set_xlabel('Unit sales')

df.groupby('cat_id')['sales'].sum().sort_values().plot.pie(ax=ax[2], startangle=45, auto
ax[2].set_title('Unit sales per category')
ax[2].set_ylabel('')

df.groupby('dept_id')['sales'].sum().sort_values().plot.barh(ax=ax[3])
ax[3].set_xlabel('Unit sales')
ax[3].set_ylabel('Departments')
ax[3].set_title('Unit sales per department')

fig.tight_layout()
```



```
In [31]: df.groupby(['d', 'cat_id'])['sales'].sum().unstack().plot(figsize=(15,5))
plt.title('Total unit sales per category since 2011')
```



```
plt.xlabel('Days')
plt.ylabel('Unit sales')
```

Out[31]: Text(0, 0.5, 'Unit sales')



```
In [32]: categories = df['cat_id'].unique()
fig, ax = plt.subplots(len(categories), 1, sharex=True, figsize=(15,8))

for i, c in enumerate(categories):
    df[df['cat_id']==c].groupby(['d', 'dept_id'], observed=True)['sales'].mean().unstack

    ax[i].set_title(f'{c} unit sales per department')
    ax[i].set_xlabel('Days')
    ax[i].set_ylabel('Average unit sales')

plt.suptitle('Average unit sales per category and department', y=0.95, fontsize='x-large')
```

Out[32]: Text(0.5, 0.95, 'Average unit sales per category and department')



Similarities across stores, departments

There are similarities both:

- in each **department**, across different states/stores, meaning that similar products behave in a similar manner,
- and in each **store**, across products belonging to different departments, meaning that there are local pattern.

In [33]: `%%capture`

```
idx = np.arange(D_PUBLIC - 60, D_PRIVATE).tolist()

# minor ticks every monday
minor_ticks = calendar.loc[idx][calendar['dayofweek'] == 0].index.tolist()
# major ticks every first day of month
major_ticks = calendar.loc[idx][calendar['dayofmonth'] == 1].index.tolist()
```

In [34]:

```
stores = df['store_id'].unique()
fig, ax = plt.subplots(10, 1, sharex=True, sharey=False, figsize=(15, 30))


for i, s in enumerate(stores):
    df[(df['store_id'] == s) & (df['d'] > D_PUBLIC-60)].groupby(['d', 'dept_id'], observ


    ax[i].get_legend().remove()
    ax[i].set_title(f'{s} unit sales per department')
    ax[i].set_xlabel('Days')
    ax[i].set_ylabel('Quantity sold')
    ax[i].set_xticks(minor_ticks, minor=True)
    ax[i].set_xticks(major_ticks, minor=False)

handles, labels = ax[-1].get_legend_handles_labels()
fig.legend(handles, labels, loc='right')
plt.suptitle('Unit sales per store (across different departments)', y=0.9, fontsize='x-1
```

Out[34]: `Text(0.5, 0.9, 'Unit sales per store (across different departments)')`

# Unit sales per store (across different departments)





```
In [35]: departments = df['dept_id'].unique()


fig, ax = plt.subplots(len(departments), sharex=True, sharey=False, figsize=(15, 18))

for i, d in enumerate(departments):
    df[(df['dept_id'] == d) & (df['d'] > D_PUBLIC-60)].groupby(['d', 'store_id'])['sales']
        .plot(linewidth=0.7, ax=ax[i])

    ax[i].get_legend().remove()
    ax[i].set_title(f'{d} sales per store')
    ax[i].set_xlabel('Days')
    ax[i].set_ylabel('Quantity sold')
    ax[i].set_xticks(minor_ticks, minor=True)
    ax[i].set_xticks(major_ticks, minor=False)

handles, labels = ax[-1].get_legend_handles_labels()
fig.legend(handles, labels, loc='right')

plt.suptitle('Unit sales per department (across different stores)', y=0.92, fontsize='x-
Out[35]: Text(0.5, 0.92, 'Unit sales per department (across different stores)')
```



## Categories across stores

By looking at whole timeseries (since 2011) per each store, it is possible to notice:

- some seasonal patterns (that will be covered in the next sessions)
- some periods which differ a lot from the general behaviour of the timeseries. These periods will be removed.

```
In [36]: fig, ax = plt.subplots(10, 1, sharex=True, figsize=(16,30))

for i, s in enumerate(df['store_id'].unique()):
    df[df['store_id']==s].groupby(['d', 'cat_id'], observed=True)['sales'].sum().unstack()

    ax[i].get_legend().remove()
```












```
ax[i].set_title(f'{s} sales per category')
ax[i].set_xlabel('Days')
ax[i].set_ylabel('Quantity sold')

handles, labels = ax[-1].get_legend_handles_labels()
fig.legend(handles, labels, loc='right')


plt.suptitle('Store unit sales per category, category', y=0.9, fontsize='x-large')
```

Out[36]: Text(0.5, 0.9, 'Store unit sales per category, category')

# Store unit sales per category, category



HOBBIES  
HOUSEHOLD  
FOODS



## Detail


- Sales in **WI\_1** are sensibly lower before day 637, and they drop down near that date
- Sales in **WI\_2** are sensibly lower before day 490, and it looks like they do not show the seasonal pattern present in the successive years

```
In [37]: fig, ax = plt.subplots(1,2, sharey=True, figsize=(15, 4))

d = 637
df[(df['store_id']=='WI_1') & (df['d']>d-50) & (df['d']<d+50)] \
    .groupby(['d', 'cat_id'], observed=True)['sales'].mean().unstack().plot(ax=ax[0])
ax[0].axvline(d-0.5, 0, 1, color='red')
ax[0].set_title('Unit sales in WI_1')

d = 490
df[(df['store_id']=='WI_2') & (df['d']>d-50) & (df['d']<d+50)] \
    .groupby(['d', 'cat_id'], observed=True)['sales'].mean().unstack().plot(ax=ax[1])
ax[1].axvline(d-0.5, 0, 1, color='red')
ax[1].set_title('Unit sales in WI_2')
```

Out[37]: Text(0.5, 1.0, 'Unit sales in WI\_2')



## Seasonality

### Year, Month

```
In [38]: categories = df['cat_id'].unique()
states = df['state_id'].unique()

fig, ax = plt.subplots(len(categories), len(states), sharex=True, sharey='row', figsize=
ax = [cell for row in ax for cell in row]

for i, c in enumerate(categories):
    data = df[(df['cat_id']==c) & (df['d']>3)]
    for j, s in enumerate(states):
        k = 3*i+j
        data[data['state_id']==s].groupby(['month', 'year'], observed=True)['sales'].mea
        ax[k].get_legend().remove()
        ax[k].set_title(f'Average {c} unit sales in {s}')
        ax[k].set_xlabel('Months')
        ax[k].set_ylabel('Average unit sales')
```



```

handles, labels = ax[-1].get_legend_handles_labels()
fig.legend(handles, labels, loc='right')

plt.suptitle('Monthly Average sales per category and state', y=0.95, fontsize='x-large')

del data

```



## Day of month

- FOODS\_2 , FOODS\_3 products are generally higher in the first half of the month.
  - This situation is enhanced in WI\_2 , WI\_3 , where there is an evident pattern.
- The day of month seems almost irrelevant for HOBBIES and HOUSEHOLD products.

```

In [39]: departments = df['dept_id'].unique()
states = df['state_id'].unique()

fig, ax = plt.subplots(len(departments), len(states), sharex=True, sharey='row', figsize=

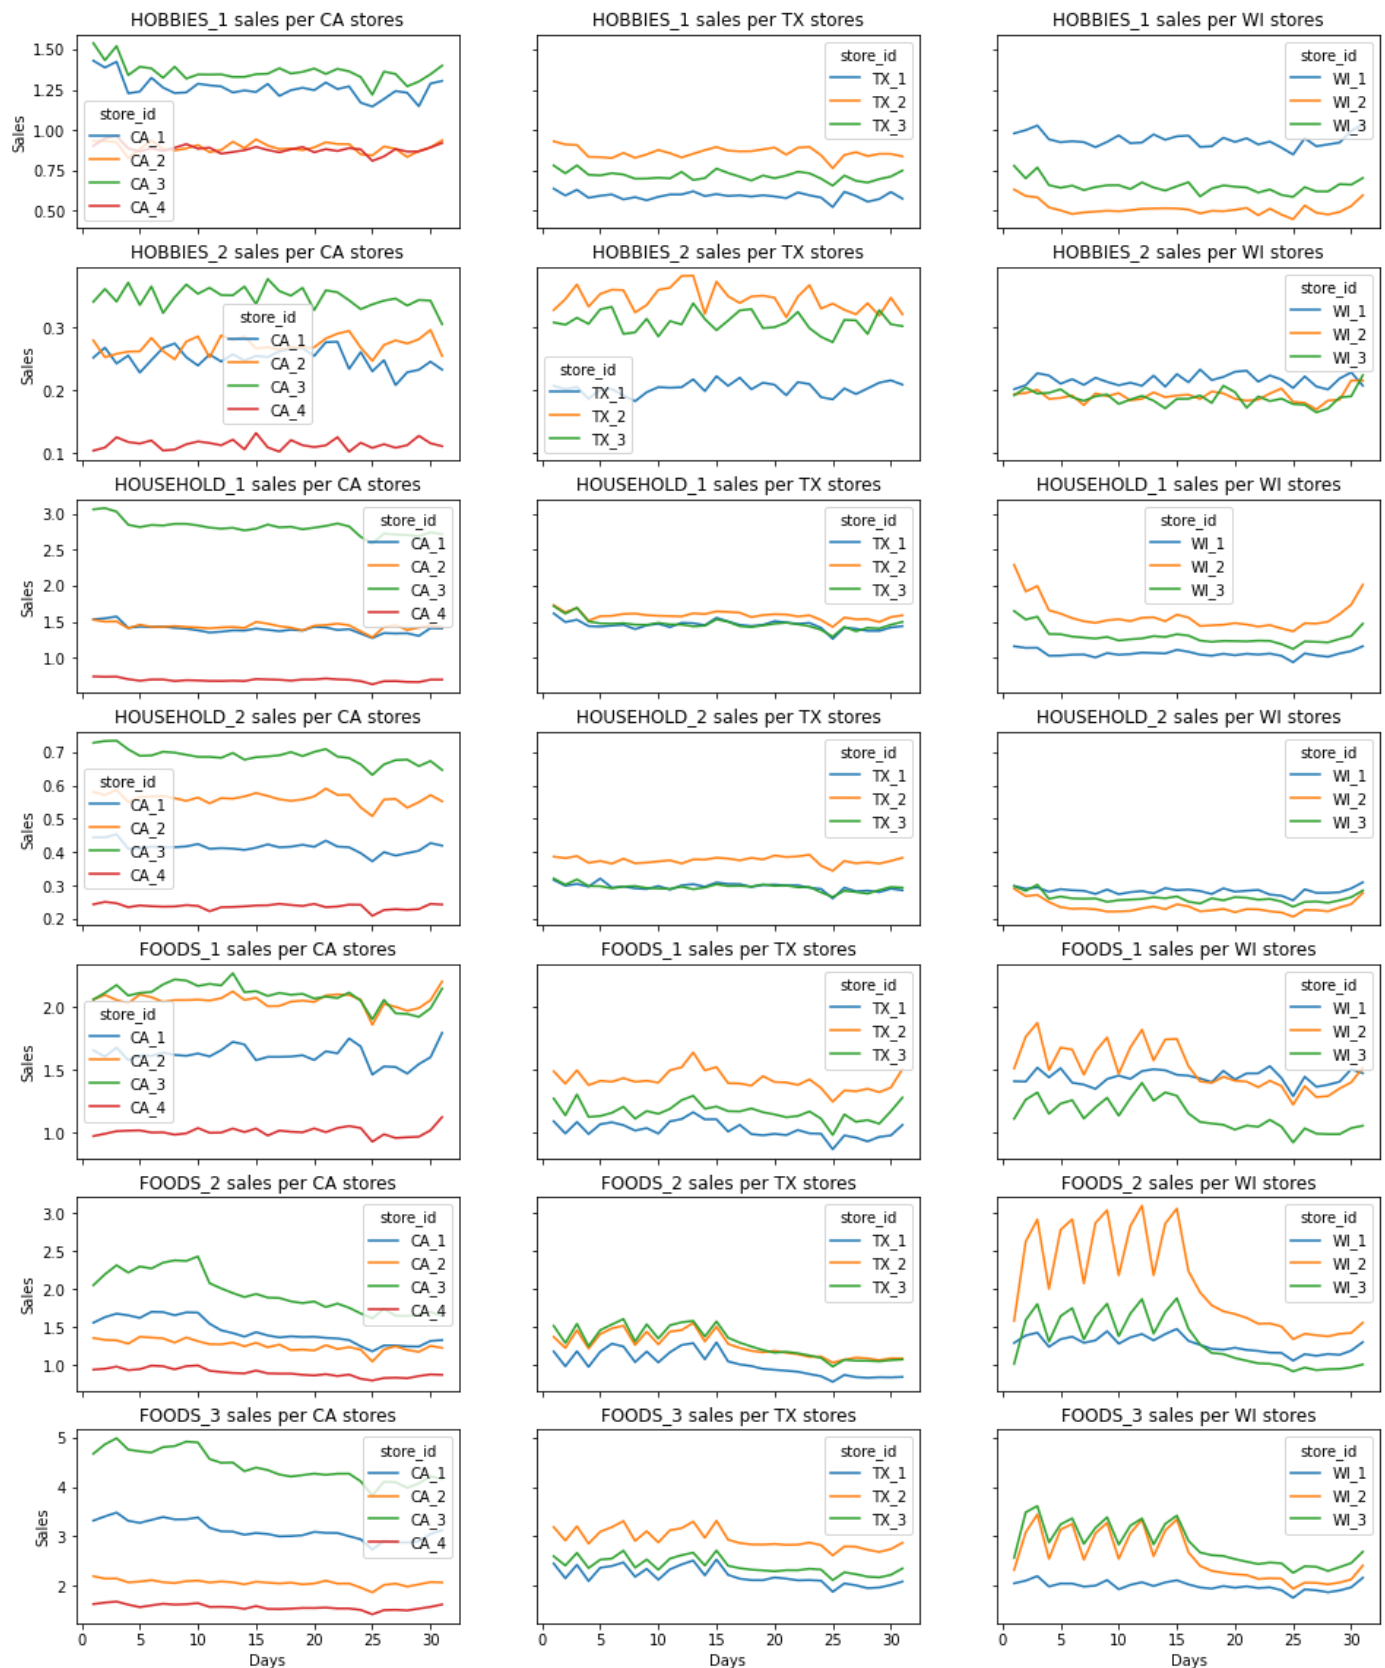
for i, d in enumerate(departments):
    data = df[df['dept_id']==d]
    for j, s in enumerate(states):
        data[data['state_id']==s].groupby(['dayofmonth', 'store_id'], observed=True)['sa

        ax[i][j].set_title(f'{d} sales per {s} stores')
        ax[i][j].set_xlabel('Days')
        ax[i][j].set_ylabel('Sales')

plt.suptitle('Daily unit sales per department and store', y=0.95, fontsize='x-large')

del data

```



## Day of week

All departments and stores have higher sales in the weekends

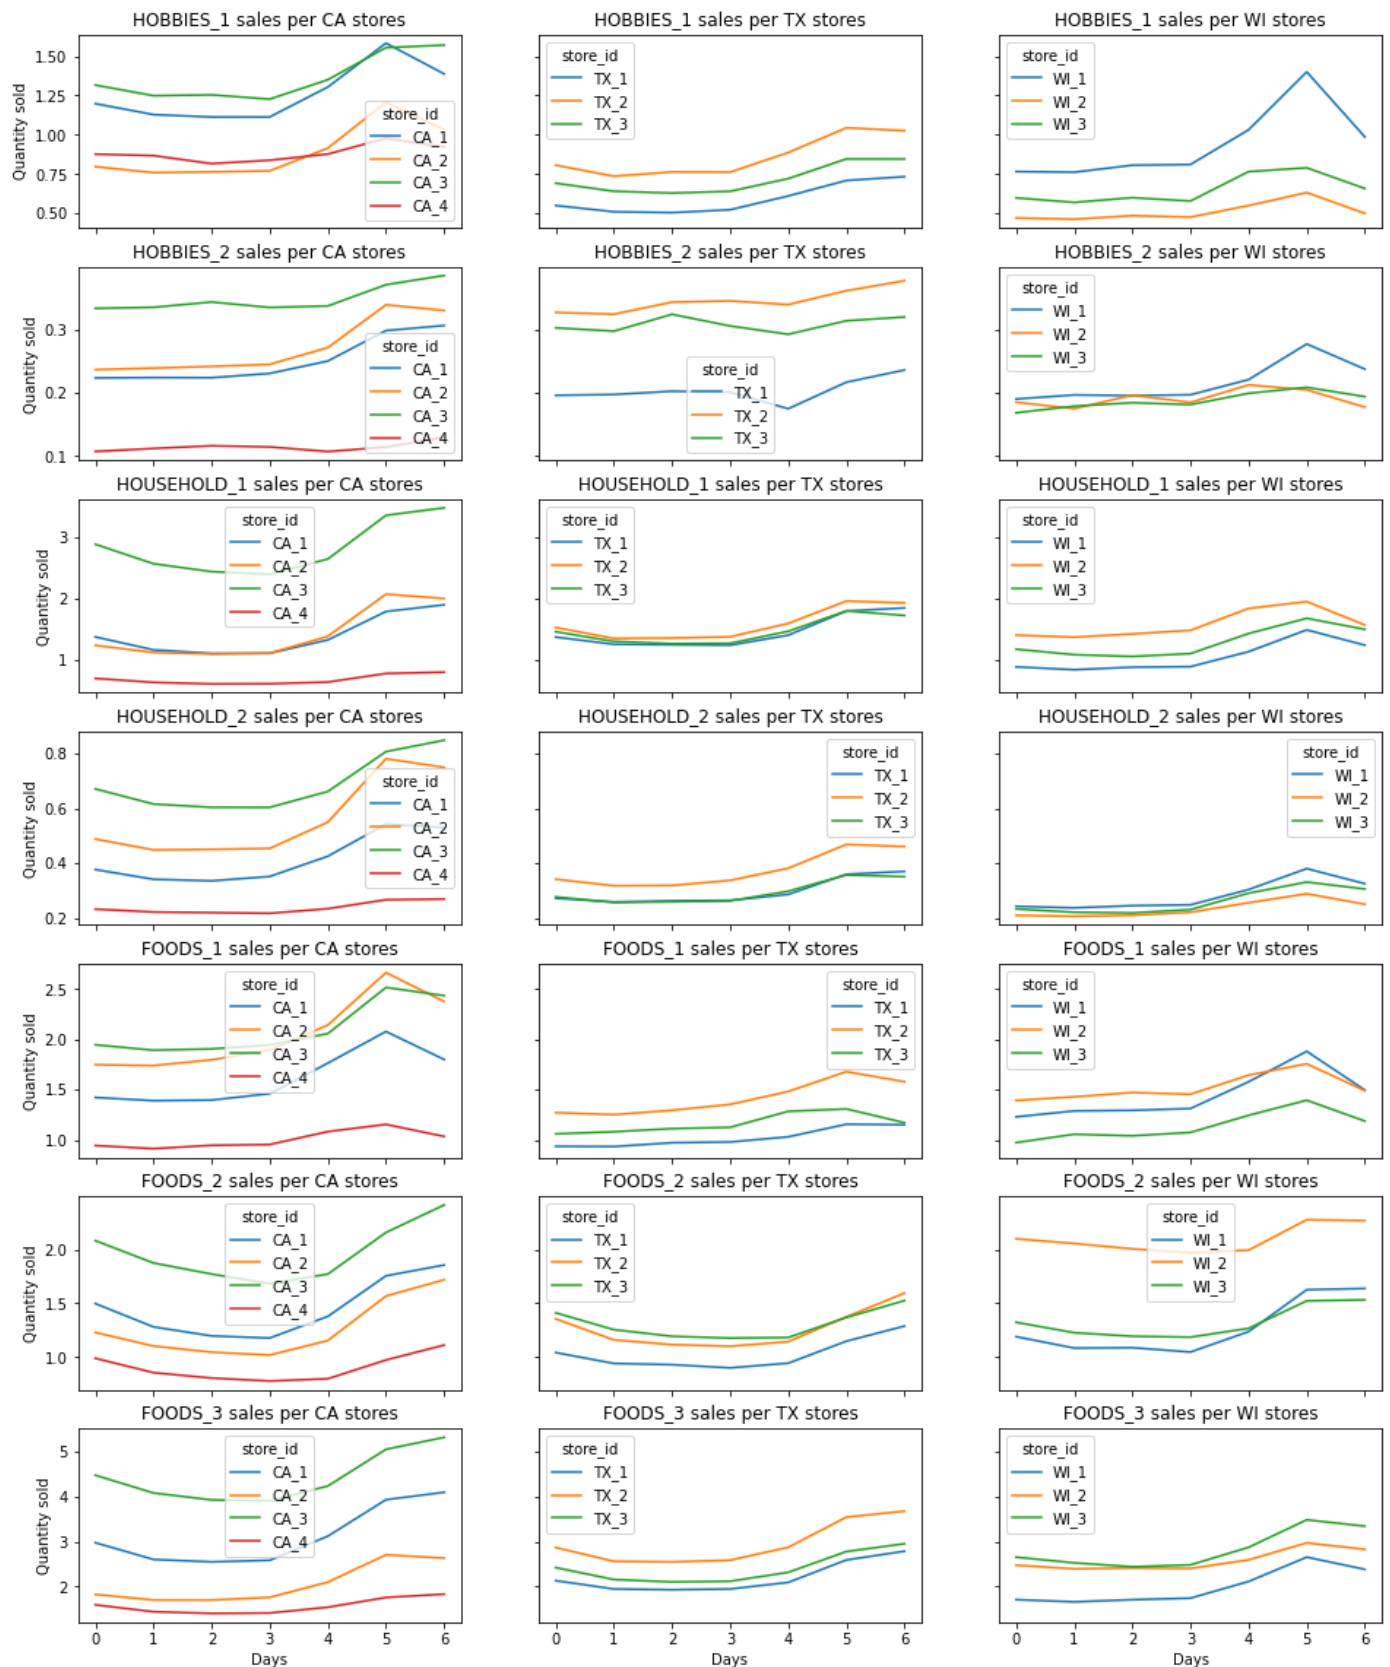
```
In [40]: departments = df['dept_id'].unique()
states = df['state_id'].unique()
```

```
fig, ax = plt.subplots(len(departments), len(states), sharex=True, sharey='row', figsize=(10, 10))

for i, d in enumerate(departments):
    data = df[df['dept_id']==d]
    for j, s in enumerate(states):
        data[data['state_id']==s].groupby(['dayofweek', 'store_id'], observed=True)['sales'].plot(
            ax=ax[i][j], title=f'{d} sales per {s} stores',
            xlabel='Days',
            ylabel='Quantity sold')

plt.suptitle('Unit sales in each day of week per store, department', y=0.92, fontsize='x-small')

del data
```



## SNAP days

- Average unit sales are usually slightly higher on SNAP days, especially for FOODS products.
- WI is the state that benefits more from SNAP promotions.

```
In [41]: states = ['CA', 'WI', 'TX']
fig, ax = plt.subplots(1, len(states), figsize=(15, 5), sharey=True)

for i, s in enumerate(states):
```


```

df[df['state_id']==s].groupby(['dept_id', 'snap'])['sales'].mean().unstack().plot.bar
ax[i].set_title(f'Average unit sales in {s}')

plt.suptitle('Average units sales on SNAP days', fontsize='x-large')

```

Out[41]: Text(0.5, 0.98, 'Average units sales on SNAP days')



## Event days

The presence of an event influences the sales. The red line shows the average unit sales on no-event days.

The behaviour is slightly different across the states, but it is possible to outline that:

- in some occasions such as SuperBowl , OrthodoxEaster and LaborDay sales are higher than the average
- On other days, such as Thanksgiving or NewYear , sales are lower, instead.
- On Christmas sales are zero since stores are closed.

```


In [42]: states = ['CA', 'WI', 'TX']
fig, ax = plt.subplots(1, len(states), figsize=(15, 8), sharey=True)

for i, s in enumerate(states):
    tmp = df[df['state_id']==s].groupby('event_name_1')['sales'].mean()
    tmp.plot.barh(ax=ax[i])
    ax[i].axvline(tmp['no_event'], color='red')
    ax[i].set_title(f'average unit sales in {s}')

plt.suptitle('Average unit sales on event days', y=0.95, fontsize='x-large')
del tmp

```

Out[42]: Text(0.5, 0.95, 'Average unit sales on event days')



## Autocorrelation

Autocorrelation is computed in order to find possible lag values that may be useful in making predictions.

For most of departments/stores, lag 7 (and its multiples) are the most relevant, but there are few exceptions:

- no/little correlation:
  - FOODS\_1 in WI\_2 , CA\_4 ;
  - HOBBIES\_1 in WI\_2 , CA\_4 ;
  - HOBBIES\_2 (all stores);
  - HOUSEHOLD\_2 in WI\_2 , CA\_4 ,
- lag 28 or 30 more relevant than lag 7:
  - FOODS\_2 in CA\_3 , TX\_1 , TX\_3 , WI\_2 , WI\_3 ;
  - FOODS\_3 in WI\_2 .

```
In [43]: from statsmodels.graphics.tsaplots import plot_acf, plot_pacf
```

```
In [44]: # average sales of product in the same department and store in different days
df2 = df.groupby(['dept_id', 'store_id', 'd'], as_index=False)['sales'].mean()
```

```
In [45]: df2.head()
```

```
Out[45]:
```

	dept_id	store_id	d	sales
0	FOODS_1	CA_1	1	4.242857
1	FOODS_1	CA_1	2	3.302326
2	FOODS_1	CA_1	3	2.431818
3	FOODS_1	CA_1	4	1.923077
4	FOODS_1	CA_1	5	1.936170

```
In [46]: departments = df2['dept_id'].unique()
stores = df2['store_id'].unique()
```

```

lags = 60

fig, ax = plt.subplots(5, 2, sharex=True, sharey=True, figsize=(15, 12))

ax = [cell for row in ax for cell in row]

cur_dept = departments[1] # change to select another department

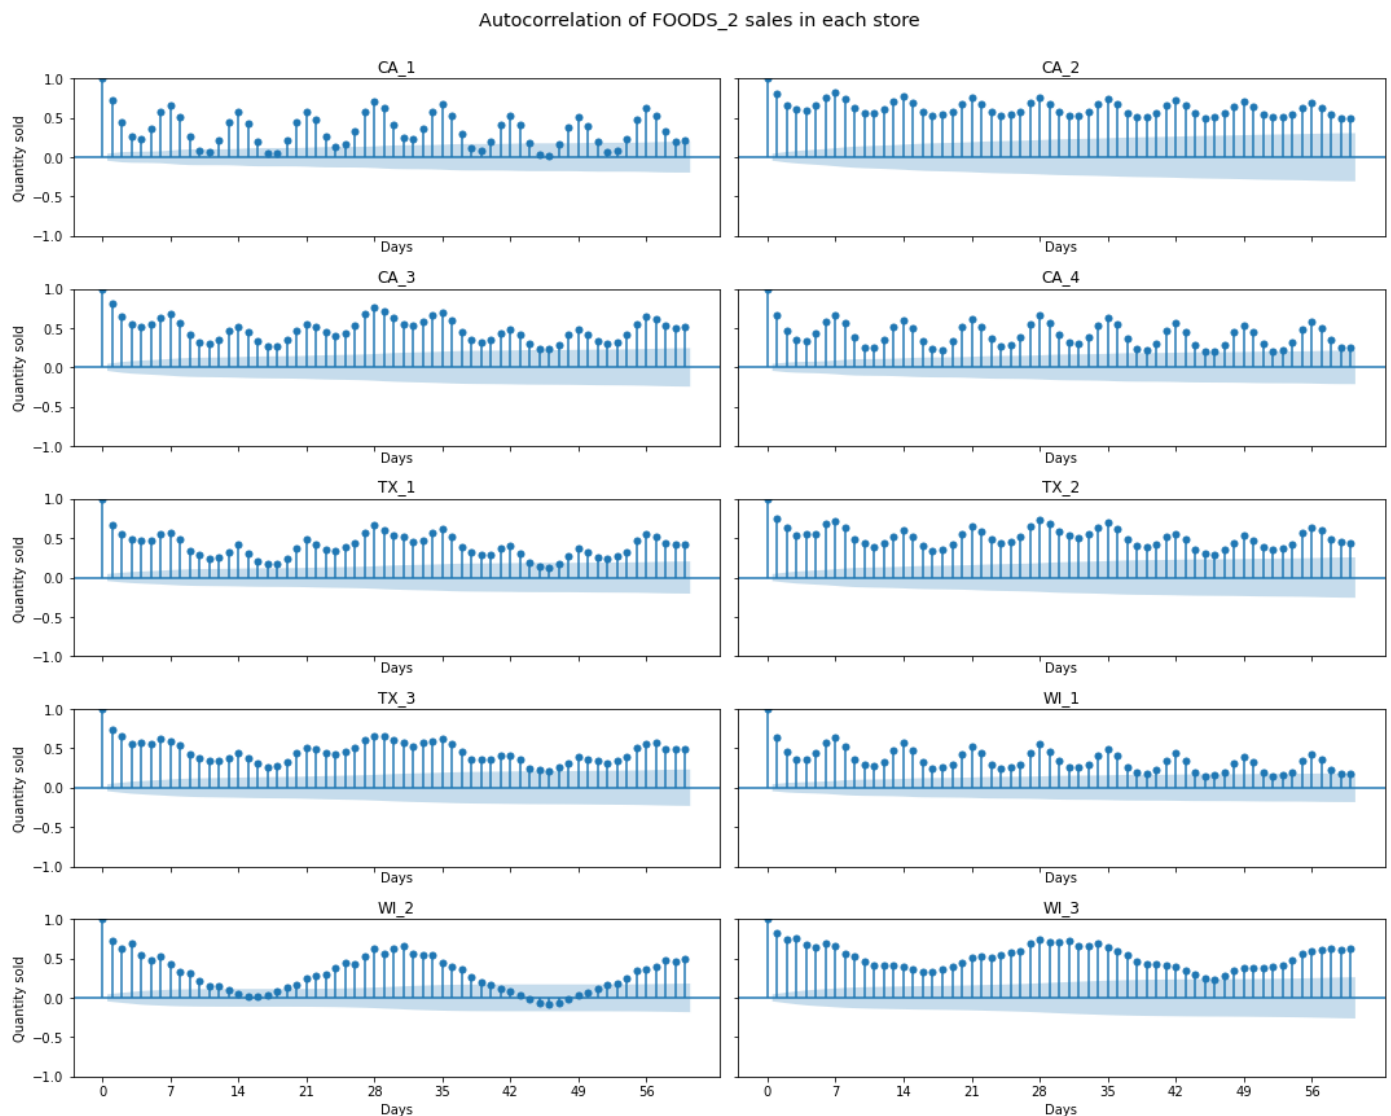
data = df2[df2['dept_id']==cur_dept]

for i, s in enumerate(stores):
    plot_acf(data[data['store_id']==s]['sales'], ax=ax[i], lags=lags)

    ax[i].set_title(f'{s}')
    ax[i].set_xlabel('Days')
    if not(i % 2):
        ax[i].set_ylabel('Quantity sold')
    ax[i].set_xticks(np.arange(0, lags, 7))

plt.suptitle(f'Autocorrelation of {cur_dept} sales in each store', y=1, fontsize='x-large')
plt.tight_layout()
del data

```



### 3) Feature engineering

Compute new features starting from available information

```
In [47]: if 'df' in globals():  
        del df
```

```
if 'df2' in globals():  
    del df2
```

```
In [49]: calendar, prices, sales = utils.read_data_pqt()  
  
prices['sell_price'] = prices['sell_price'].astype('float16')
```

```
In [55]: %%time  
  
# change to select another state  
cur_state = 'WI' # 'CA', 'TX', 'WI'  
  
df = sales[(sales['state_id']==cur_state)]  
df.reset_index(drop=True, inplace=True)  
  
df = preprocessing.add_days(df)  
df = preprocessing.remove_leading_zeros(df, calendar)  
  
if cur_state == 'WI':  
    df = df[(df['store_id']!='WI_1') | (df['d']>636)]  
    df = df[(df['store_id']!='WI_2') | (df['d']>489)]  
    df.reset_index(inplace=True, drop=True)  
  
fe.compute_and_save_products(df, cur_state)  
fe.compute_and_save_prices(df, prices, cur_state)  
fe.compute_and_save_sales(df)  
  
Saving products...  
Computing prices...  
Computing sales...  
CPU times: total: 3min 25s  
Wall time: 3min 25s
```

## 4) Model Implementation

```
In [ ]: D_MAX = df['d'].max()  
D_MIN = df['d'].min()  
D_PRIVATE = D_MAX - 28 + 1 # 1942 - 1969  
D_PUBLIC = D_PRIVATE - 28 # 1914 - 1941  
  
print(f'Training range:\t {D_MIN} - {D_PRIVATE - 1}')  
print(f'Evaluation range: {D_PRIVATE} - {D_MAX}')
```

```
In [57]: if 'prices' in globals():  
        del prices  
if 'calendar' in globals():  
    del calendar  
if 'sales' in globals():  
    del sales  
if 'df' in globals():  
    del df
```

## Load dataframe related to a specific store

```
In [377... cur_store = 'WI_3' # change to select a different store
```

```
In [378... df = utils.load_df(cur_store)
```




```
In [379... df.drop(columns=['store_id', 'state_id', 'wm_yr_wk'], inplace=True)
```

## Correlation

Correlation of each column wrt the label ( sales )

```
In [380... corr = df.corrwith(df['sales']).sort_values()
corr.plot.barh(figsize=(5, 10))
plt.grid(True)
plt.title('Correlation of each column wrt the target (sales)')

del corr
```



```
In [381... df.head()
```

```
Out[381]:
```

		id	d	sales	rolling_mean_dayofweek	rolling_mean_7_28	rolling_mean_28_28	rolling_mean_28_90
0	HOBBIES_1_002_WI_3	338		0.0	1.000000	1.500000	0.678711	0.500000
1	HOBBIES_1_004_WI_3	338		0.0	2.333984	2.535156	2.357422	2.000000
2	HOBBIES_1_005_WI_3	338		2.0	1.666992	2.892578	2.322266	2.100000
3	HOBBIES_1_008_WI_3	338		0.0	3.500000	4.355469	5.035156	3.600000
4	HOBBIES_1_009_WI_3	338		2.0	1.666992	1.428711	1.142578	0.900000

5 rows × 39 columns

```
In [382]: df.info()

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 4312966 entries, 0 to 4312965
Data columns (total 39 columns):
 #   Column                                Dtype
---  -
 0   id                                    category
 1   d                                    int16
 2   sales                                float16
 3   rolling_mean_dayofweek               float16
 4   rolling_mean_7_28                    float16
 5   rolling_mean_28_28                   float16
 6   rolling_mean_28_90                   float16
 7   rolling_mean_28_180                  float16
 8   rolling_mean_28_365                  float16
 9   rolling_mean_dayofmonth               float16
10   avg_item_store                        float16
11   std_item_store                        float16
12   avg_item_state                        float16
13   avg_dept_store                        float16
14   avg_dept_state                        float16
15   avg_snap                              float16
16   avg_event                             float16
17   avg_weekofmonth                       float16
18   n_items                               int16
19   item_id                              category
20   dept_id                              category
21   cat_id                               category
22   release                              int16
23   weekofmonth                          int8
24   weekofyear                           int8
25   dayofweek                             int8
26   dayofmonth                           int8
27   dayofyear                             int16
28   month                                int8
29   year                                 int16
30   event_name_1                         category
31   event_type_1                         category
32   snap_WI                              int8
33   sell_price                           float16
34   price_last_week                      float16
35   price_next_week                      float16
36   price_same_dept                      float16
37   price_mean                           float16
38   price_std                            float16
dtypes: category(6), float16(22), int16(5), int8(6)
memory usage: 281.0 MB
```

```
In [385]: x_cols = df.columns.tolist()
# label is removed
```

```
x_cols.remove('sales')

# remove id/item_id for better generalization
x_cols.remove('id')
x_cols.remove('item_id')
# remove other columns not used by the model
x_cols.remove('avg_dept_state')
x_cols.remove('year')
x_cols.remove('event_type_1')
```

In [386... x\_cols

```
Out[386]: ['d',
'rolling_mean_dayofweek',
'rolling_mean_7_28',
'rolling_mean_28_28',
'rolling_mean_28_90',
'rolling_mean_28_180',
'rolling_mean_28_365',
'rolling_mean_dayofmonth',
'avg_item_store',
'std_item_store',
'avg_item_state',
'avg_dept_store',
'avg_snap',
'avg_event',
'avg_weekofmonth',
'n_items',
'dept_id',
'cat_id',
'release',
'weekofmonth',
'weekofyear',
'dayofweek',
'dayofmonth',
'dayofyear',
'month',
'event_name_1',
'snap_WI',
'sell_price',
'price_last_week',
'price_next_week',
'price_same_dept',
'price_mean',
'price_std']
```

## LGBM model

In [387... `import lightgbm as lgbm`  
`lgbm.__version__`

Out[387]: '3.3.2'

## Public and private data

Public data is used as validation, private data as evaluation (once submitted)

In [390... `mask_valid = (df['d']>=D_PUBLIC) & (df['d']<D_PRIVATE)`  
`mask_eval = df['d']>=D_PRIVATE`

```
# public data
x_valid = df[mask_valid][x_cols]
y_valid = df[mask_valid]['sales']
```

```
# private data
x_eval = df[mask_eval][x_cols]
```

```
In [391... def cv_training(params, df, x_cols, target='sales', cv=3, dim_cv=28, early_stopping=100,
    """
    Train a LGBM model with cross validation
    @param params: LGBM parameters
    @param df: input dataframe
    @param x_cols: input columns
    @param target: label column
    @param cv: number of cross validation periods
    @param dim_cv: length of each cv period
    @param early_stopping: number of epochs for early stopping
    @param d_max: last day used in validation
    @param lr: starting learning rate
    @param min_lr: minimum learning rate
    @return LGBM boosters
    """

    # start/end days of the cross validation
    d_val = [d_max - dim_cv*i for i in range(cv+1)]
    d_val.reverse()
    print('CV ranges:')
    for i in range(cv):
        print(f'CV {i+1}: {d_val[i]} - {d_val[i+1]-1}')

    # train indices
    tr_idx = [df[df['d']==d].index[0] for d in d_val[:-1]]
    tr_indices = [pd.RangeIndex(0, i) for i in tr_idx]
    # validation indices
    val_idx = [df[df['d']==d].index[0] for d in d_val[1:]]
    val_indices = [pd.RangeIndex(tr_idx[i], val_idx[i]) for i in range(cv)]

    ds = lgbm.Dataset(df[x_cols], df[target])

    # early stopping, adaptive learning rate, log evaluation every 20 steps
    cb = [lgbm.early_stopping(early_stopping),
          lgbm.reset_parameter(learning_rate=lambda iter: max(min_lr, lr*(0.99**iter))),
          lgbm.log_evaluation(period=20, show_stdv=False)]

    booster = lgbm.cv(params,
                      folds=list(zip(tr_indices, val_indices)),
                      train_set=ds,
                      num_boost_round=2000,
                      shuffle=False,
                      callbacks=cb,
                      return_cvbooster=True,
                      eval_train_metric=True,
                      seed=0)

    return booster
```

```
In [392... # parameters for training the odel
lr = 0.1
params = {'boosting': 'gbdt',
          'num_leaves': 2**5-1,
          'objective': 'tweedie',
          'tweedie_variance_power': 1.1, #1.1, 1.3., 1.5
          'metric': 'rmse',
          'learning_rate': lr,
          'bagging_freq': 10, 'bagging_fraction': 0.25,
          'force_col_wise': True,
          'deterministic': True
        }
```

```

In [393... boost = cv_training(params, df, x_cols, d_max=D_PRIVATE, lr=lr)

CV ranges:
CV 1: 1858 - 1885
CV 2: 1886 - 1913
CV 3: 1914 - 1941
[LightGBM] [Info] Total Bins 6005
[LightGBM] [Info] Number of data points in the train set: 3971478, number of used features: 33
[LightGBM] [Info] Total Bins 6005
[LightGBM] [Info] Number of data points in the train set: 4056850, number of used features: 33
[LightGBM] [Info] Total Bins 6005
[LightGBM] [Info] Number of data points in the train set: 4142222, number of used features: 33
[LightGBM] [Info] Start training from score 0.252069
[LightGBM] [Info] Start training from score 0.252210
[LightGBM] [Info] Start training from score 0.252575
Training until validation scores don't improve for 100 rounds
[20]    cv_agg's train rmse: 2.57496    cv_agg's valid rmse: 2.15739
[40]    cv_agg's train rmse: 2.43478    cv_agg's valid rmse: 2.04282
[60]    cv_agg's train rmse: 2.39848    cv_agg's valid rmse: 2.02069
[80]    cv_agg's train rmse: 2.37389    cv_agg's valid rmse: 2.02407
[100]   cv_agg's train rmse: 2.35354    cv_agg's valid rmse: 2.02136
[120]   cv_agg's train rmse: 2.33343    cv_agg's valid rmse: 2.04788
[140]   cv_agg's train rmse: 2.31513    cv_agg's valid rmse: 2.04443
[160]   cv_agg's train rmse: 2.29742    cv_agg's valid rmse: 2.03964
Early stopping, best iteration is:
[70]    cv_agg's train rmse: 2.3846 + 0.0066803 cv_agg's valid rmse: 2.01546 + 0.0287035

```

## Model selection

For each of the 3 models returned, RMSE on validation is computed. The model with the smallest error is maintained.

```

In [394... models = boost['cvbooster']

In [395... y_valid_pred = models.predict(x_valid)

errors = [mse(y_valid, y, squared=False) for y in y_valid_pred]

print('RMSE of each cvbooster on public data:')
for i, err in enumerate(errors):
    print(f'CVBooster {i}:\t{err: .3f}')

RMSE of each cvbooster on public data:
CVBooster 0:    2.091
CVBooster 1:    2.005
CVBooster 2:    1.994

```

```

In [396... np.argmin(errors)

```

Out[396]: 2

```

In [397... model = models.boosters[np.argmin(errors)]

```

## Feature importance

Have a look at the importance of each input feature in the model


```

In [400... plt.figure(figsize=(5,10))

```

```
pd.Series(model.feature_importance(), model.feature_name()).sort_values().plot.barh()
```

Out[400]: <AxesSubplot:>



In [401... `model.feature_importance()`

Out[401]: `array([322, 241, 496, 318, 145, 102, 85, 126, 120, 76, 45, 4, 69, 503, 198, 137, 118, 31, 33, 19, 123, 218, 208, 296, 2, 164, 68, 118, 124, 75, 151, 256, 109])`

In [402... `dict(zip(model.feature_name(), model.feature_importance()))`

```
Out[402]: {'d': 322,
           'rolling_mean_dayofweek': 241,
           'rolling_mean_7_28': 496,
           'rolling_mean_28_28': 318,
           'rolling_mean_28_90': 145,
           'rolling_mean_28_180': 102,
           'rolling_mean_28_365': 85,
           'rolling_mean_dayofmonth': 126,
           'avg_item_store': 120,
           'std_item_store': 76,
           'avg_item_state': 45,
           'avg_dept_store': 4,
           'avg_snap': 69,
           'avg_event': 503,
           'avg_weekofmonth': 198,
           'n_items': 137,
           'dept_id': 118,
           'cat_id': 31,
           'release': 33,
           'weekofmonth': 19,
           'weekofyear': 123,
           'dayofweek': 218,
           'dayofmonth': 208,
           'dayofyear': 296,
           'month': 2,
           'event_name_1': 164,
           'snap_WI': 68,
           'sell_price': 118,
           'price_last_week': 124,
           'price_next_week': 75,
           'price_same_dept': 151,
           'price_mean': 256,
           'price_std': 109}
```

```
In [404... # save the model
           #utils.save_model(model, f'm5_{cur_store}.txt')
```

```
In [405... # load the model
           #model = utils.load_model(store=cur_store)
```

```
In [406... cur_store
```

```
Out[406]: 'WI_3'
```

## 5) Predictions

```
In [407... y_valid_pred = utils.predict(df, D_PUBLIC, model, x_cols)
           print(f'RMSE: {mse(y_valid, y_valid_pred, squared=False):.3f}')

           RMSE: 2.049
```

```
In [408... y_eval_pred = utils.predict(df, D_PRIVATE, model, x_cols)
```

## 6) Submission

### Partial submission

Compute public and private submission for a **single** store

```
In [409... valid_submission = utils.partial_submission(df.loc[mask_valid, ['id', 'd']], y_valid_pre
```

```

eval_submission = utils.partial_submission(df.loc[mask_eval, ['id', 'd']], y_eval_pred,

submission = pd.concat([valid_submission, eval_submission])

dst = '../partial_submissions'
if not os.path.exists(dst):
    os.makedirs(dst)

submission.to_csv(f'{dst}/m5_{cur_store}.csv', index=False)

```

## Final submission

Compute final submission for **all** stores

```
In [410]: final_submission = utils.final_submission()
```

## 7) Evaluation

Compute WRMSSE on public data and decompose it for all aggregation levels

```
In [1]: # extract data if not done already
        #utils.extract_data()
```

```
In [6]: calendar, prices, sales = utils.read_data_csv()

        sales = sales.sort_values(by='id').reset_index(drop=True)

        train_df = sales.iloc[:, :-28]
        valid_df = sales.iloc[:, -28:]
```

```
In [7]: evaluator = wrmsse_evaluator.WRMSSEEvaluator(train_df, valid_df, calendar, prices)

        100%|██████████| 12/12 [00:21<00:00, 1.75s/it]
```

```
In [9]: final_submission = pd.read_csv('../final_submission/m5_final_submission.csv')
        valid_submission = final_submission[final_submission['id'].str[-10:]=='validation']
        valid_submission = valid_submission.sort_values(by='id').reset_index(drop=True)
```


```
In [18]: wrmsse = evaluator.score(valid_submission.iloc[:, -28:].values)
         print(f'WRMSSE on public data: {wrmsse:.5f}')

        WRMSSE on public data: 0.60637
```

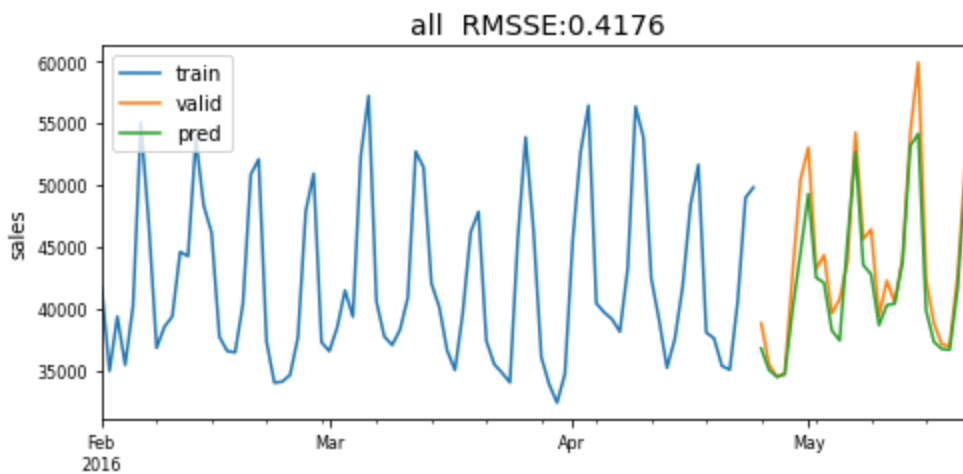
```
In [19]: wrmsse_evaluator.create_dashboard(evaluator)
```




## WRMSSE by Level



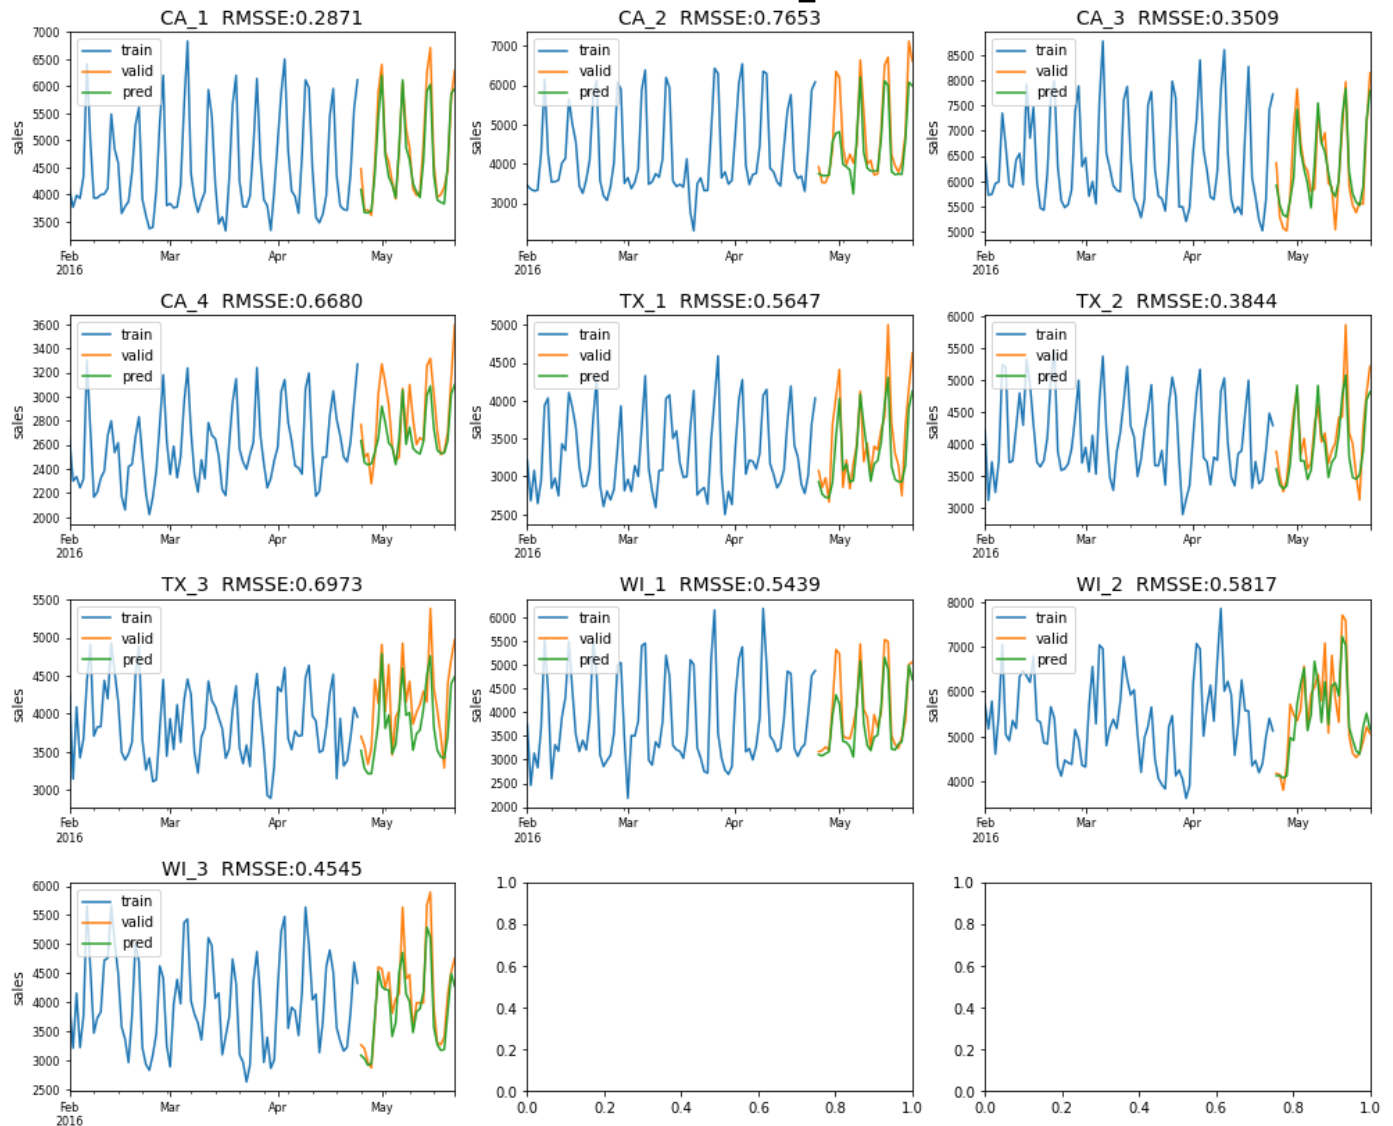
### Level 1: all\_id




### Level 2: state\_id












### Level 3: store\_id



### Level 4: cat\_id



## Level 5: dept\_id



In [ ]: